



Getting Started with Geospatial Data in MongoDB

Buzz Moschetti
Enterprise Architect
buzz.moschetti@mongodb.com
[@buzzmoschetti](#)

Agenda

- What is MongoDB?
- What does “geospatial capabilities” mean?
- GeoJSON
- Combining GeoJSON with non-geo data
- APIs and Use Cases
- Comparison to OGC (Open Geospatial Consortium)
- Indexing
- Using Geo Capabilities for non-Geo Things
- Esri and shapefiles

MongoDB:

The Post-Relational General Purpose Database



Fully Featured
High Performance
Scalable

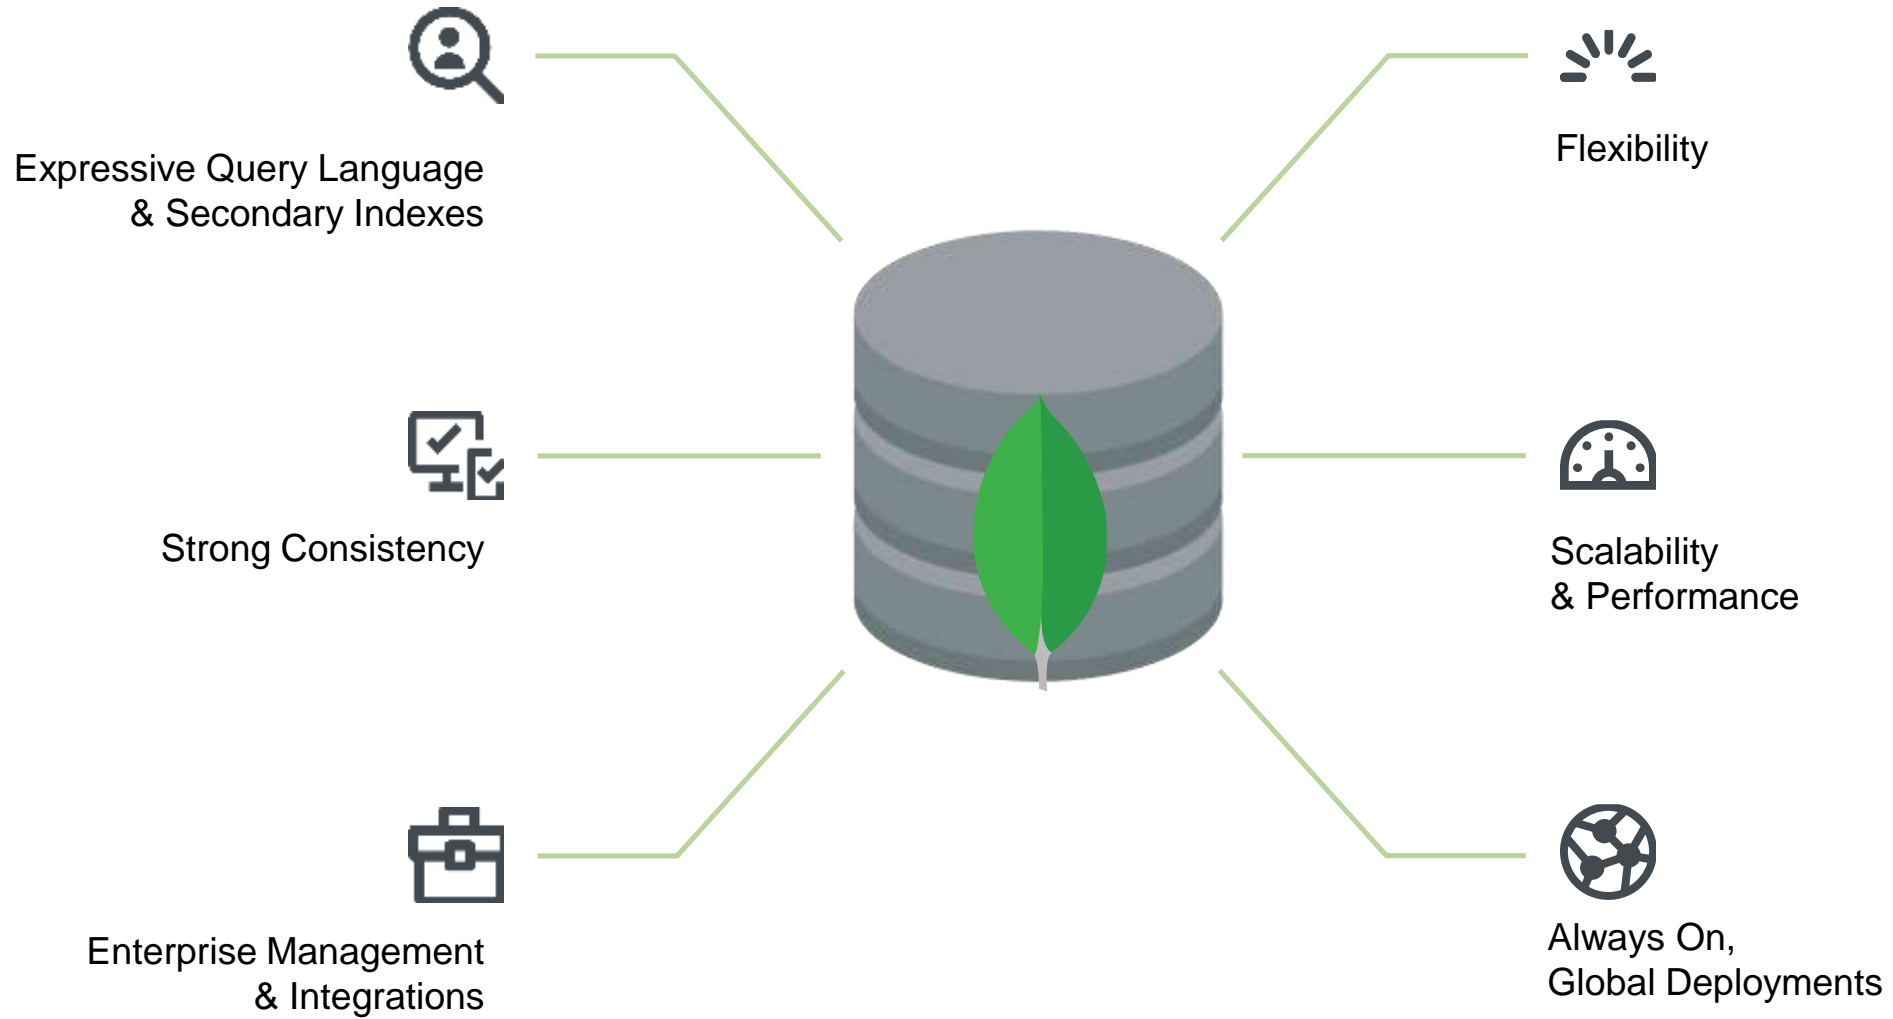
```
{  
  name: "John Smith",  
  pfxs: ["Dr.", "Mr."],  
  address: "10 3rd St.",  
  phone: {  
    home: 1234567890,  
    mobile: 1234568138 }  
}
```

Document
Data Model



Open-
Source

Nexus Architecture




MongoDB Company Overview



~800 employees



2500+ customers



Offices in NY & Palo Alto and
across EMEA, and APAC




Over \$311 million in funding

What is “Geo”?

At least **4 levels** of capability




The Geo Stack



Efficiently **store, query, and index** arbitrary points, lines and **polygons** in the DB


The Geo Stack



Platform for data analysis of **peer data** (trades/house value/population/sales/widgets) grouped by **geo data**

Efficiently **store, query, and index** arbitrary points, lines and polygons in the DB

The Geo Stack




Graphical rendering of geo shapes on a **map**

Platform for data analysis of **peer data** (trades/house value/population/sales/widgets) grouped by geo data

Efficiently **store, query, and index** arbitrary points, lines and polygons in the DB

The Geo Stack



Application(s) to browse and manipulate **all** the data

Graphical rendering of geo shapes on a map

Platform for data analysis of **peer data** (trades/house value/population/sales/widgets) grouped by geo data

Efficiently **store, query, and index** arbitrary points, lines and polygons in the DB

Important: Sometimes there is NO Map

- Geo stack must support geo functions **WITHOUT** a Map
- Offline reporting
 - “Nightly fleet management report”
 - “Distributor loss by assigned area”
- Compute/analytical processing
 - Dynamic polygon generation
 - Weather catastrophe simulation
 - Other geo-filtering as input to analytics


MongoDB: The Data “Base”

Application(s) to browse and visualize

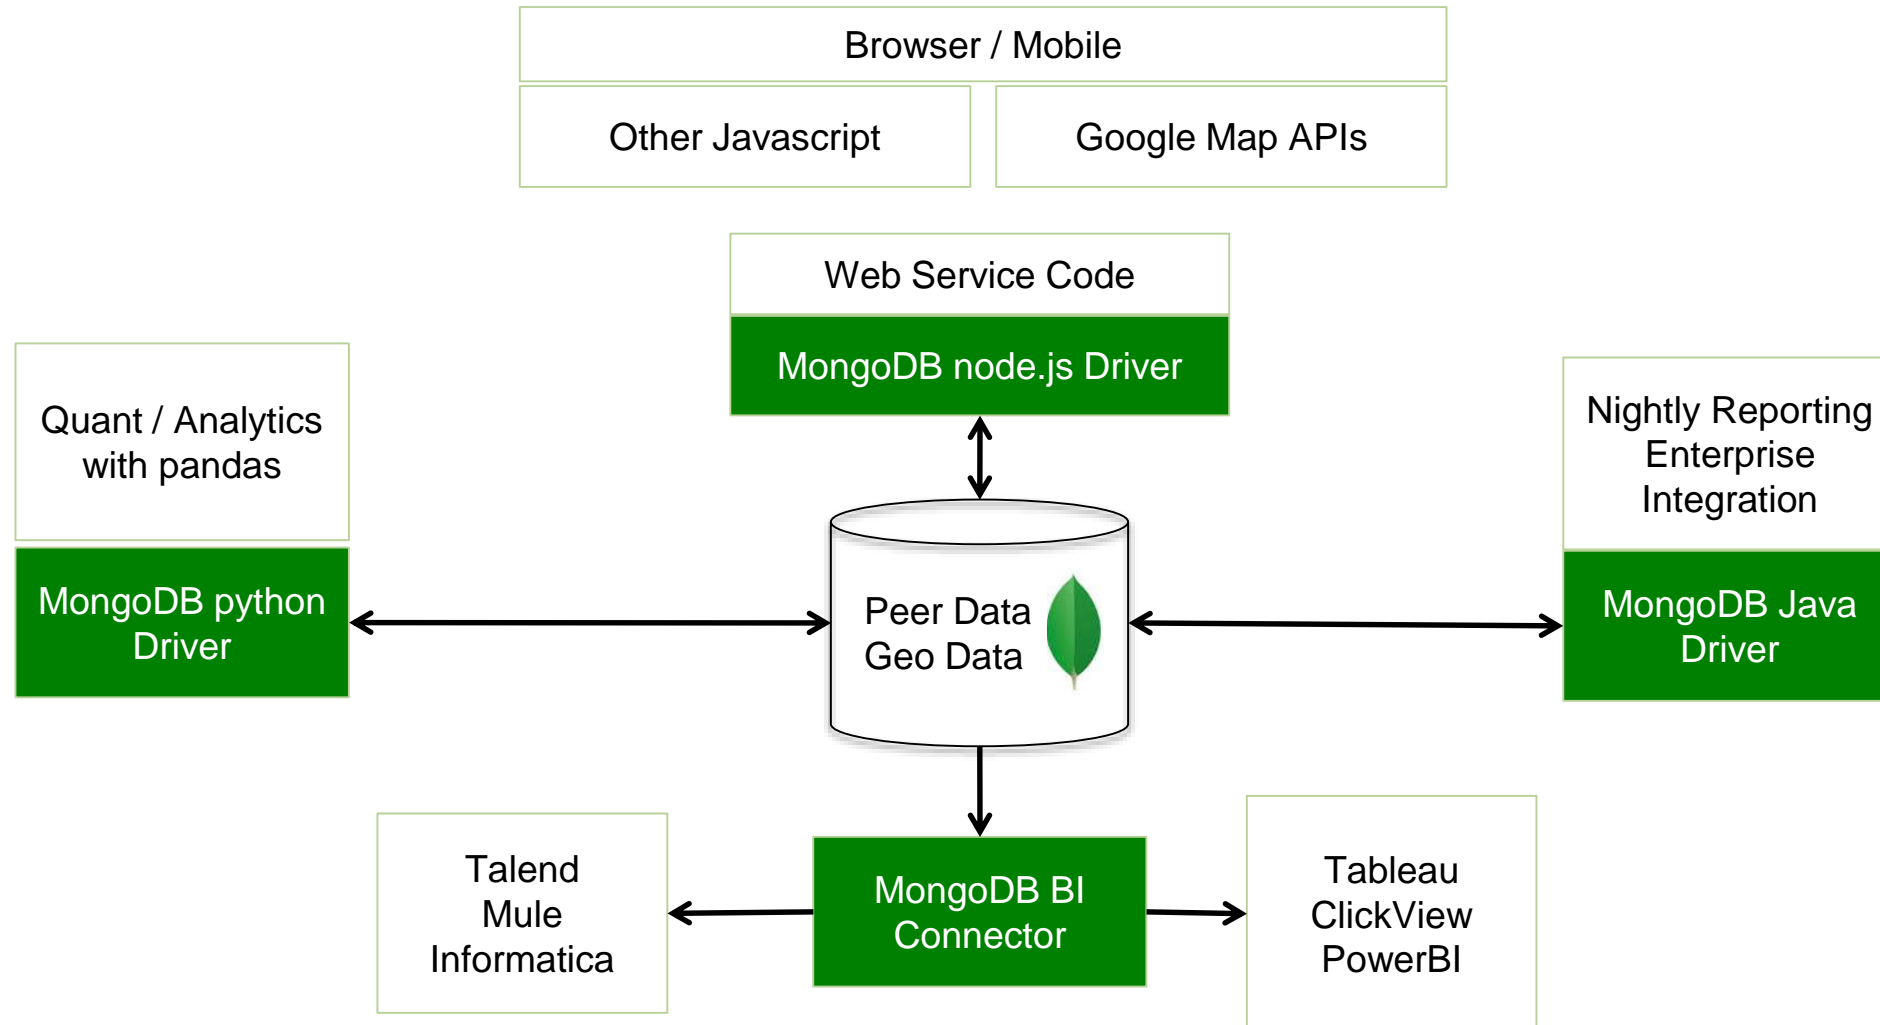
Graphical rendering of geo shapes on a map

Platform for data analysis of **peer data**
(trades/value/population/sales/widgets) grouped by
geo data

Efficiently **store, query, and index** arbitrary points,
lines, and polygons in the DB



One Persistor for All Applications & Use Cases



Balanced Reporting

- Most other NoSQL DBs **do not** have this capability
- Oracle, Postgres, MySQL, SQLServer **do** offer it and subscribe to Open GeoS Consortium (OGC) standards

MongoDB data model is the major difference

MongoDB: Simple, parse-free, type-correct APIs and data to manipulate and interrogate geo shapes

a.k.a. arrays (of arrays (of arrays))

OpenGIS: Piles of “ST_” functions:

http://postgis.net/docs/reference.html#Geometry_Accessors

```
SELECT ST_MakePolygon(  
  ST_GeomFromText(  
    'LINESTRING(75.15 29.53,77 29,77.6 29.5, 75.15 29.53)')  
  );
```

Data & APIs



Legacy: 2D points

```
{  
  name: {f: "Buzz", l: "Moschetti"},  
  favoritePets: [ "dog", "cat" ],  
  house: [ -95.12345, 43.23423 ]  
}
```

Better: GeoJSON

```
{  
  name: {f: "Buzz", l: "Moschetti"},  
  favoritePets: [ "dog", "cat" ],  
  house: {  
    type: "Point",  
    coordinates: [ -95.12345, 43.23423 ]  
  }  
}
```

Better: GeoJSON

```
{
  name: "Superfund132",
  location: {
    type: "Polygon",
    coordinates: [
      [ [-95.12345, 43.2342], [-95.12456, 43.2351], ... ]
      [ [-92.8381, 43.75], ... ] // "hole"
    ]
  }
}
```

The GeoJSON Family

```
{  
  type: "Point", "MultiPoint", "LineString", "MultiLineString", "Polygon",  
  "MultiPolygon"  
  coordinates: [ specific to type ]  
}
```

```
{  
  type: "GeometryCollection"  
  geometries: [  
    { type: (one of above),  
      coordinates: [ . . . ]  
    }  
  ]  
}
```

NO COMPUTED SHAPES
(Circle, Arc, Box, etc.)

We use the **WGS84** standard:
<http://spatialreference.org/ref/epsg/4326/>

MongoDB Data Types are Geo-friendly

```
var poly = [  
  [ [-95.12345, 43.2342], [-95.12345, 43.2351],  
    [-95.12211, 43.2351], [-95.12211, 43.2342],  
    [-95.12345, 43.2342] // close the loop!  
  ]  
];  
  
db.myCollection.insert(  
  {name: {f: "Buzz", l: "Moschetti"},  
   favoritePets: ["dog", "cat"],  
   geo: { type: "Polygon", coordinates: poly }  
}));
```


... even with Java

```
Document doc = new Document();  
doc.put("name", "Superfund132");
```

```
List ring = new ArrayList();  
addPoint(ring, -95.12345, 43.2342);  
addPoint(ring, -95.12345, 43.2351);  
addPoint(ring, -95.12211, 43.2351);  
addPoint(ring, -95.12211, 43.2342);  
addPoint(ring, -95.12345, 43.2342);
```

```
List poly = new ArrayList();  
poly.add(ring);  
Map mm = new HashMap();  
mm.put("type", "Polygon");  
mm.put("coordinates", poly);  
doc.put("geo", mm);
```

```
coll.insertOne(doc);
```



```
static void addPoint(List ll,  
                    double lng,  
                    double lat) {  
    List pt = new ArrayList();  
    pt.add(lng);  
    pt.add(lat);  
    ll.add(pt);  
}
```

All Types Are Preserved Correctly

```
Document doc = coll.find().first();  
recursiveWalk(doc);
```

```
name: java.lang.String: Superfund132  
geo: com.mongodb.BasicDBObject  
  type: java.lang.String: Polygon  
  coordinates: com.mongodb.BasicDBList  
    0: com.mongodb.BasicDBList  
      0: com.mongodb.BasicDBList  
        0: java.lang.Double: -95.12345  
        1: java.lang.Double: 43.2342  
      1: com.mongodb.BasicDBList  
        0: java.lang.Double: -95.12345  
        1: java.lang.Double: 43.2351  
    2: com.mongodb.BasicDBList  
      0: java.lang.Double: -95.12211  
      1: java.lang.Double: 43.2351
```

**MongoDB Geo
is just doubles
and Lists!**

Comparison to “Good” PostGIS

```
import org.postgis.PGgeometry; // extended from org.postgresql.util.PGobject

((org.postgresql.Connection) conn).addDataType("geometry", "org.postgis.PGgeometry")

String sql = "select geom from someTable";
ResultSet r = stmt.executeQuery(sql);
while( r.next() ) {
    PGgeometry geom = (PGgeometry)r.getObject(1);
    if( geom.getType() = Geometry.POLYGON ) {
        Polygon pl = (Polygon)geom.getGeometry();
        for( int r = 0; r < pl.numRings(); r++) {
            LinearRing rng = pl.getRing(r);
            . . .
        }
    }
}
```

**Beware of
bespoke types
and
dependencies!**

Comparison to most OpenGLS

```
String sql = "select ST_AsText(geom) from someTable";
ResultSet r = stmt.executeQuery(sql);
while( r.next() ) {
    String wkt = r.getString(1);

    // wkt is "POLYGON((0 0,0 1,1 1,1 0,0 0))"

    // http://en.wikipedia.org/wiki/Well-known\_text

    // Now we have to parse the string into
    // an array of array of doubles.
    // Don't want to introduce a 3rd party dependency...
    // So . . . We write our own parser.
}
```

Checkpoint

We have data in and out of the DB using basic operations (insert and find)

Now we need to make it performant!

Indexing

```
collection.createIndex({loc:"2d"})
```

When to use:

- Your database has legacy location data from MongoDB 2.2 or earlier
- You do not intend to store **any** location data as GeoJSON objects
- “Special Use Cases” e.g. arbitrary two numeric dimension indexing

```
collection.createIndex({loc:"2dsphere"})
```

When to use:

- Supports all GeoJSON objects **and legacy [x,y] pairs**

```
collection.createIndex({loc:"geoHaystack"})
```

When to use:

- Special small area flat (planar) lookup optimization

Indexing

```
collection.createIndex({loc:"2d"})
```

When to use:

- Your database has legacy location data from MongoDB 2.2 or earlier
- You do not intend to store **any** location data as GeoJSON objects
- “Special Use Cases” e.g. arbitrary two numeric dimension indexing

```
collection.createIndex({loc:"2dsphere"})
```

When to use:

- Supports all GeoJSON objects **and legacy [x,y] pairs**

```
collection.createIndex({loc:"geoHaystack"})
```

When to use:

- Special small area flat (planar) lookup optimization



find()/\$match and Indexing

Operator	Geometry Arg Type	2d	2dsphere
\$geoWithin	\$box,\$center,\$polygon	Y	N
	\$geometry: { type, coordinates }	N	Y
	\$centerSphere: [[x,y], radians]	Y	Y
\$geoIntersects	\$geometry only	N	Y
\$near,\$nearSphere	[x,y]	R	-
(output sorted by distance)	\$geometry: {type, coordinates}	-	R
	+ \$minDistance	N	Y
	+ \$maxDistance	Y	Y

Y = will assist

N = will not assist

R = REQUIRED

Syntax helper:

```
find("loc":{$geoWithin: {$box: [ [x0,y0], [x1,y2] ]}});
```

```
find("loc":{$geoWithin: {$geometry: { type: "Polygon", coordinates: [ .... ] }}} );
```

Aggregation Framework: \$geoNear

	Option	2D	2dsphere
\$geoNear (output sorted by distance)	near: { type: "Point", coordinates }	-	R - and spherical:true
	near: [x, y]	R (or)	R
	query: { expression INCL geo find() on previous page EXCEPT \$near }	N	N

Important Considerations:

1. You can only use \$geoNear as the first stage of a pipeline.
2. You must include the distanceField option.
3. The collection **must have only one geospatial index**: one 2d index or one 2dsphere index.
4. You do not need to specify which field in the documents hold the coordinate pair or point.
Because \$geoNear requires that the collection have a single geospatial index, \$geoNear implicitly uses the indexed field.

Y = will assist

N = will not assist

R = REQUIRED

Use Cases



Case #1: Find Things in a Given Area + More

- Docs contain Points (or possibly “small” polygons)
- \$geoWithin

```
db.site.aggregate([
  { $match: { "loc": { $geoWithin: { $geometry:
    { type: "Polygon", coordinates: [ coords ] }}}
    , "portfolio_id": portid
    , "insuredValue": { $gt: 1000000 }
    , "insuredDate": { $gt: new ISODate(„2016-01-01“) } }
  , { $bucket: { groupBy: „$insuredValue“,
    boundaries: [ 1000000, 2000000, 5000000, 10000000,
      20000000, Infinity ] } }
  . . .
```


Case #2: Find Things in an Area Stored in DB

- Get the shape from the “shapes” collection via query:

```
db.shapes.findOne({predicate},{theShape:1});
```

- Turn around and query the target collection, e.g. buildingSites with shape:

```
db.buildingSites.find({loc:{$geoWithin: theShape}})
```

Case #3: Find Things Closest to where I am

```
db.buildingSites.aggregate([{$geoNear: { point ... }}]);
```

- Results returned **already in sorted order by closeness**

Case #3.5: Find Things Closest to where I am but within some bounds

- `db.buildingSites.aggregate([
 { $geoNear: {
 query: { "loc": { $geoWithin:
 { $centerSphere: ... } } }
 }])`

(or)

```
    query: { "loc": { $geoWithin: { $geometry:  
GeoJSON } } }  
  } ])
```

When the Database isn't enough



When the Database isn't enough

- VERY fast intersection/within for many objects given probes at high velocity (10000s/sec).
- Geo manipulation: unions, deltas, layering
- Dynamic/programmatic geo construction
- Advanced features: smoothing, simplifiers, centroids, ...

You Need Three Things

- Basic geo objects
- Geo operators like intersects, within, etc.
- Algos and smoothers, etc.

Don't forget
The Geo Stack!

com.vividsolutions.jts


```
Map m = (Map) dbo.get("loc"); // get a GeoJSON object from MongoDB
List coords = (List) m.get("coordinates");
List outerRing = (List) coords.get(0); // polygon is array of array of pts

CoordinateSequence pseq = new CoordinateGeoJSONSequence(outerRing, true);
LinearRing outerlr = new LinearRing(pseq, gf);
int numHoles = coords.size() - 1; // -1 to skip outer ring;
LinearRing[] holes = null;
if(numHoles > 0) {
    holes = new LinearRing[numHoles];
    for(int k = 0; k < numHoles; k++) {
        List innerRing = (List) coords.get(k+1); // +1 adj for outer ring
        CoordinateSequence psx = new CoordinateGeoJSONSequence(innerRing, true);
        holes[k] = new LinearRing(psx, gf);
    }
}
Polygon poly1 = new Polygon(outerlr, holes, gf); // ok if holes was null
Point pt1 = new Point(pseq2, gf);
boolean a = pt1.intersects(poly1);
Geometry simplified = TopologyPreservingSimplifier.simplify(poly1, tolerance);
```

The Ecosystem

- OpenGeo runs over MongoDB!
<http://suite.opengeo.org/docs/latest/dataadmin/mongodb/index.html>
- **BoundlessGeo**: Commerical support for OpenGeo over MongoDB
 - * Provides top 2 tiers (viz, analysis)
 - * <https://boundlessgeo.com>

Geo Capabilities beyond “Simple Geo”



Geo as Date Range

```
{  
  who: 'john'  
  where: 'mongodb'  
  what: 'lightning talk'  
  start:   ISODate("2016-06-30T15:00:00")  
  end:     ISODate("2016-06-30T15:05:00")  
}
```

What events were happening at 15:03?

```
collection.find({  
  start : { $lte:ISODate("2016-06-30T15:05:03")},  
  end:    { $gte:ISODate("2016-06-30T15:05:03")}  
})
```

Geo as Date Range

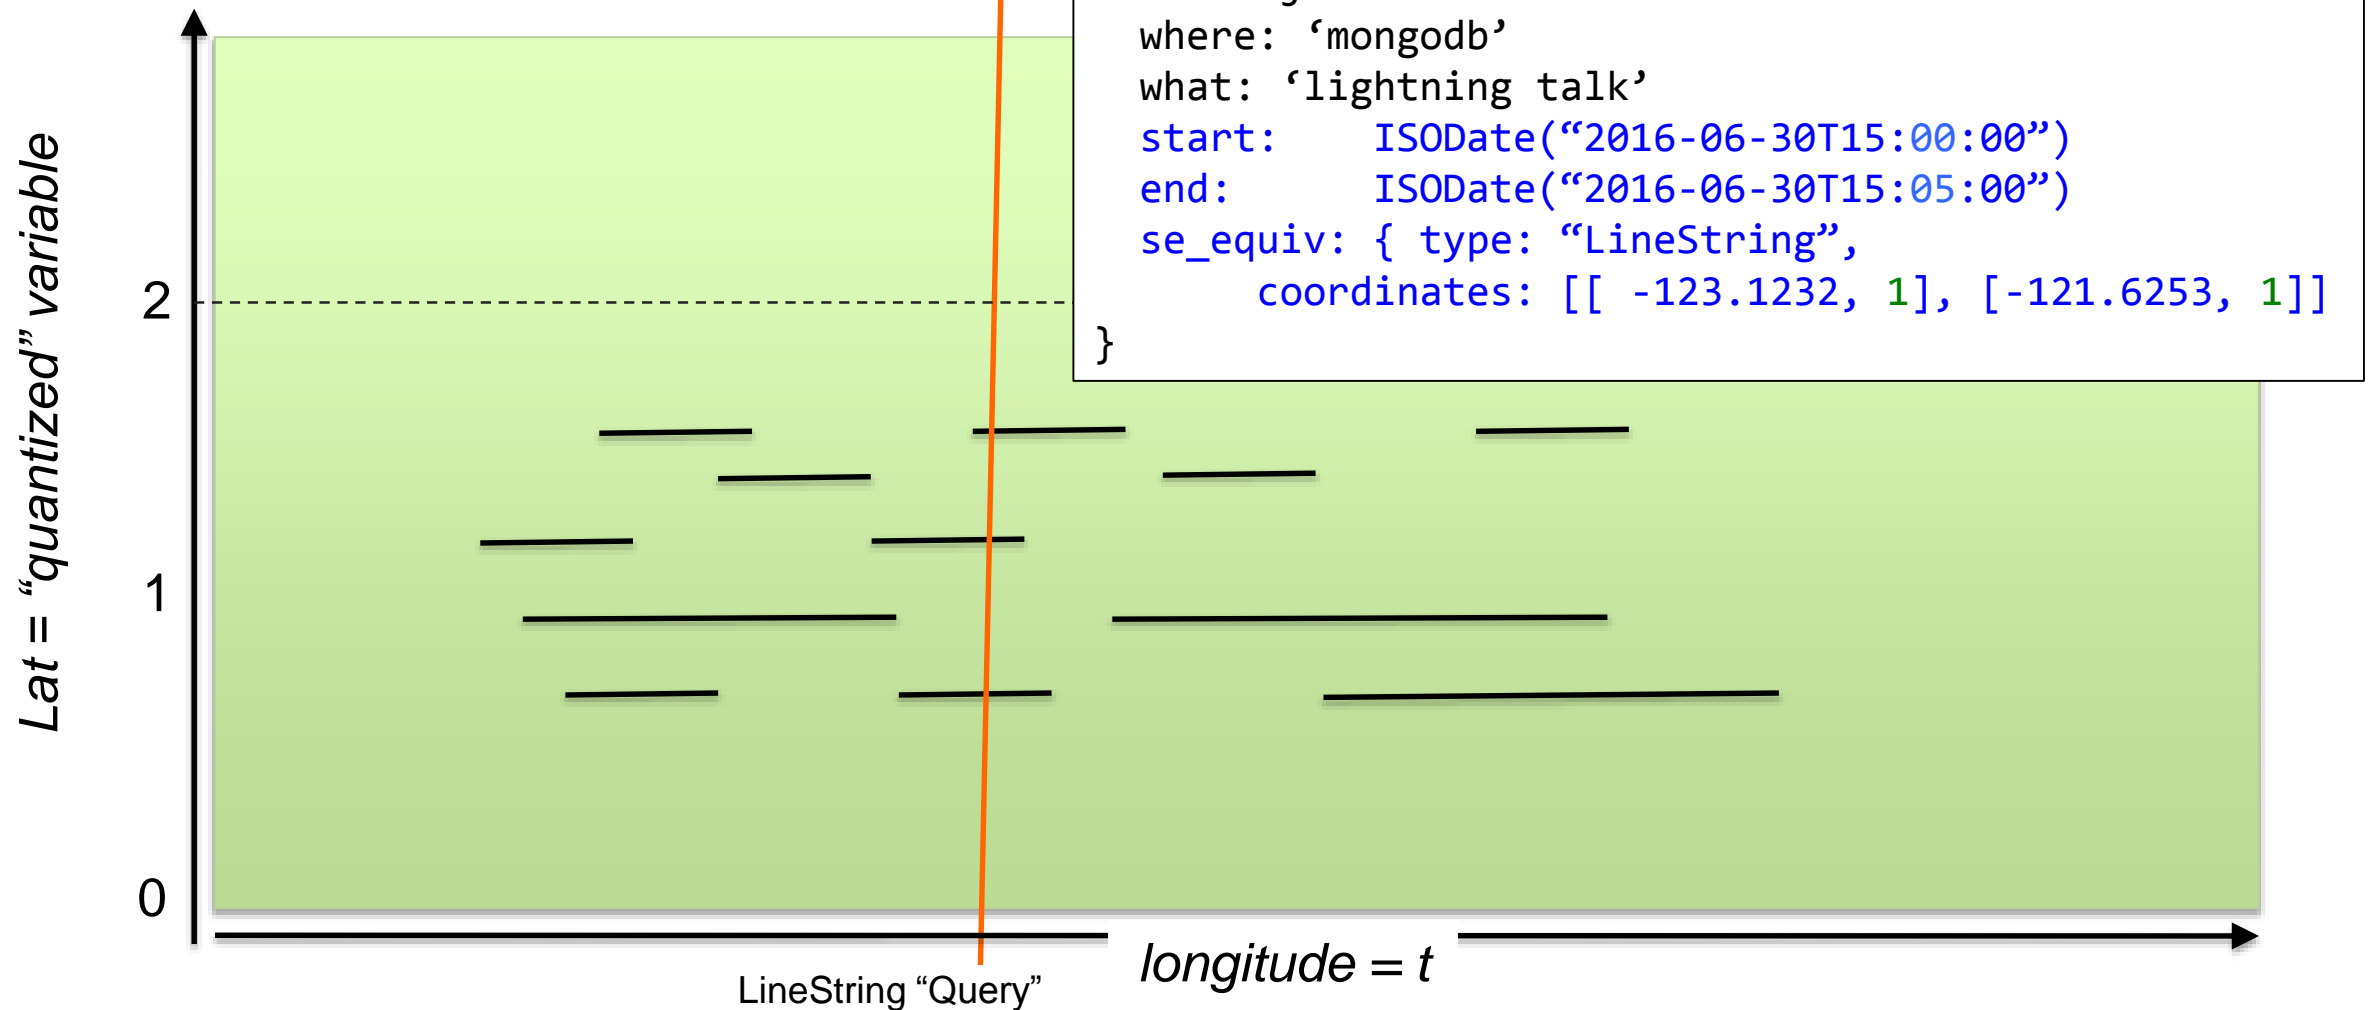
- Ranges on 2 attributes – Two BTree walks (intersection)
- Assuming time can be anywhere in range of records, index walk is average 50% of index
- Test: Macbook Pro, i5, 16GB RAM, data fits in WT Cache easily. Warmed up. Average of 100 runs

694ms /query using index

487ms /query – COLLSCAN!



(StartDate,EndDate) → Range Type using Geo



Over 10X performance improvement!

```
start2 = (((start / yearsecs) - 45) * 90) - 90
end2 = (((end / yearsecs) - 45) * 90) - 90
event = { type: "LineString", coordinates: [ [ start2, 1 ], [end2, 1 ] ] }
```

```
// dx = is the LineString "query"
```

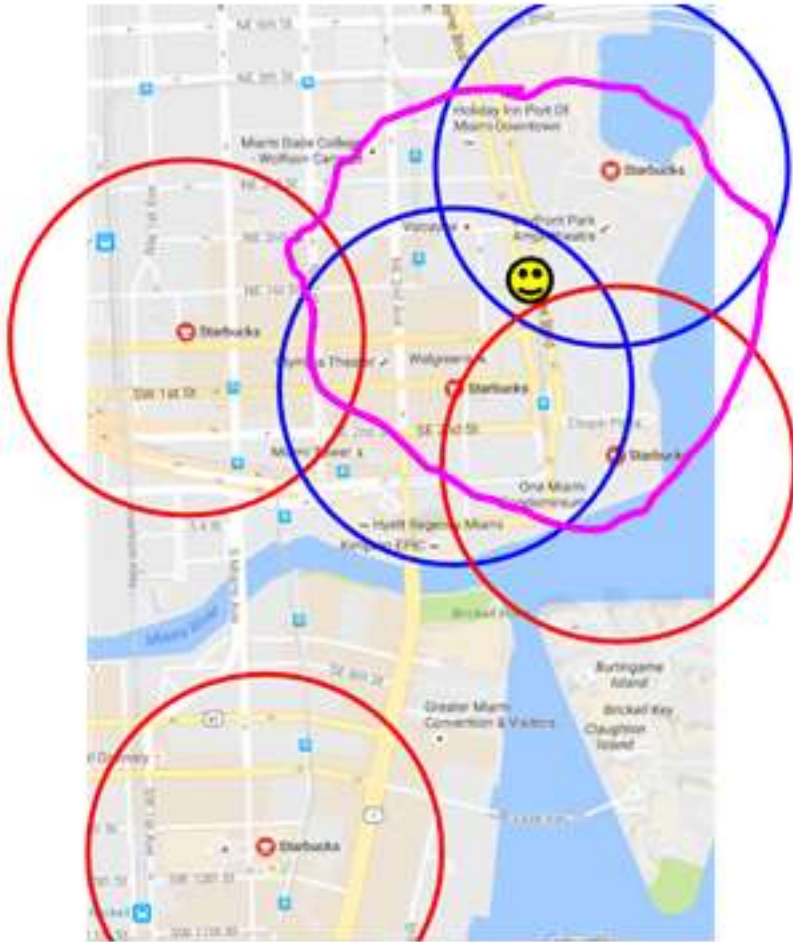
```
query = {
  g: {
    $geoIntersects: {
      $geometry: { type: "LineString",
        coordinates: [ [ dx, 0 ], [dx, 2 ] ] }
    }
  }
}
```

45ms!



Mr. Smiley

```
// Assume Mr. Smiley has these params:  
var mslng = -90.00;  
var mslat = 42.00;  
var msr = 0.005; // ~1500 ft radius around him
```



```
db.participatingVendors.aggregate([
```

```
// Stage 1: The Bounding Box  
{ $match: { "loc": { $geoWithin: { $geometry:  
    { type: "Polygon", coordinates: mapBoundingBox }}}  
}}
```

```
// Stage 2: Compute distance from Mr. Smiley to the points: Pythagorean theorem:  
, { $addFields: { dist: { $sqrt: { $add: [  
    { $pow: [ { $subtract: [ { $arrayElemAt: [ "$loc.coordinates", 0 ] }, mslng ] }, 2 ] },  
    { $pow: [ { $subtract: [ { $arrayElemAt: [ "$loc.coordinates", 1 ] }, mslat ] }, 2 ] },  
    ] } } }  
}}
```

```
// Stage 3: If the distance between points is LTE the sum of the radii, then  
// Mr. Smiley's circle intersects that of the participant:  
// Project 0 (no touch) or 1 (touch)  
, { $addFields: { "touch": { $cond: [  
    { $lte: [ { $add: [ "$promoRadius", msr ] }, "$dist" ] }, 0, 1  
    ] } } }
```

```
, { $match: { "touch": 1 } }  
]);
```

The Pusher

Don't forget
\$slice!

```
var pts = [ [-74.01,40.70], [-73.99,40.71], . . .];

db.foo.insert({_id:1,
               loc: { type:"LineString", coordinates: [ pts[0], pts[1] ]}});

// Push points onto LineString to "extend it" in an index optimized way!
for(i = 2; i < pts.length; i++) {
    db.foo.update({_id:1},{$push: {"loc.coordinates": pts[i]}});

    // Perform other functions, e.g.
    c=db.foo.find({loc: {$geoIntersects:
                        {$geometry: { type: "Polygon", coordinates: ... } } } });
}
```

Perimeter of Simple Polygon

```
> db.foo.insert({_id:1, "poly": [ [0,0], [2,12], [4,0], [2,5], [0,0] ] });
> db.foo.insert({_id:2, "poly": [ [2,2], [5,8], [6,0], [3,1], [2,2] ] });

> db.foo.aggregate([
  {$project: {"conv": {$map: { input: "$poly", as: "z", in: {
    x: {$arrayElemAt: ["$$z",0]}, y: {$arrayElemAt: ["$$z",1]}
    ,len: {$literal: 0} }}}}},
  {$addFields: {first: {$arrayElemAt: [ "$conv", 0 ]} }},
  {$project: {"qqq":
    {$reduce: { input: "$conv", initialValue: "$first", in: {
      x: "$$this.x", y: "$$this.y"
      ,len: {$add: ["$$value.len", // len = oldlen + newLen
        {$sqrt: {$add: [
          {$pow: [ {$subtract: ["$$value.x", "$$this.x"]}, 2]}
          , {$pow: [ {$subtract: ["$$value.y", "$$this.y"]}, 2]}
        ] } } ] } } }
    },
  {$project: {"len": "$qqq.len"}}

  { "_id" : 1, "len" : 35.10137973546188 }
  { "_id" : 2, "len" : 19.346952903339393 }
```


Geospatial = 2D Numeric Indexable Space

Find all branches close to my location:

```
target = [ someLatitude, someLongitude ];
radians = 10 / 3963.2;      // 10 miles
db.coll.find({"location": { $geoWithin :
                           { $center : [ target, radians ] }}}));
```

Find nearest investments on efficient frontier:

```
target = [ risk, reward ];
closeness = someFunction(risk, reward);
db.coll.find({"investmentValue": { $geoWithin :
                                   { $center : [ target, closeness ] }}}));
```

Basic Tips & Tricks

- We use [long,lat], not [lat,long] like Google Maps
- Use 2dsphere for geo; avoid legacy \$box/\$circle/\$polygon
- Use 2d for true 2d numeric hacks
- 5 digits beyond decimal is accurate to 1.1m:
 - `var loc = [-92.12345, 42.56789] // FINE`
 - `var loc = [-92.123459261145, 42.567891378902] // ABSURD`
- \$geoWithin and \$geoIntersects do not REQUIRE index
- Remember to close loops (it's GeoJSON!)

esri-related Tips & Tricks

- Shapefiles are **everywhere**; google shapefile <whatever>
- Crack shapefiles to GeoJSON with python [pyshp](#) module

```
import shapefile
import sys
from json import dumps

reader = shapefile.Reader(sys.argv[1])

field_names = [field[0] for field in reader.fields[1:] ]
buffer = []

for sr in reader.shapeRecords():
    buffer.append(dict(geometry=sr.shape.__geo_interface__,
                      properties=dict(zip(field_names, sr.record))))

sys.stdout.write(dumps({"type": "FeatureCollection", "features": buffer},
indent=2) + "\n")
```



Q & A



Thank You!

Buzz Moschetti
Enterprise Architect
buzz.moschetti@mongodb.com
[@buzzmoschetti](#)

Agenda

- What does geospatial capabilities mean?
- The "levels": DB with geo types, rendering, analytics
- MongoDB brings together geo AND non-geo data
- Geo Data model
- GeoJSON
- Combining GeoJSON with non-geo data
- APIs and Use Cases
- Looking up things contained in a polygon
- Finding things near a point
- Summary of geo ops e.g. \$center
- \$geoNear and the agg framework
- The power of the document model and MongoDB APIs
- Arrays and rich shapes as first class types
- Comparison to Postgres (PostGIS)

\

Indexing

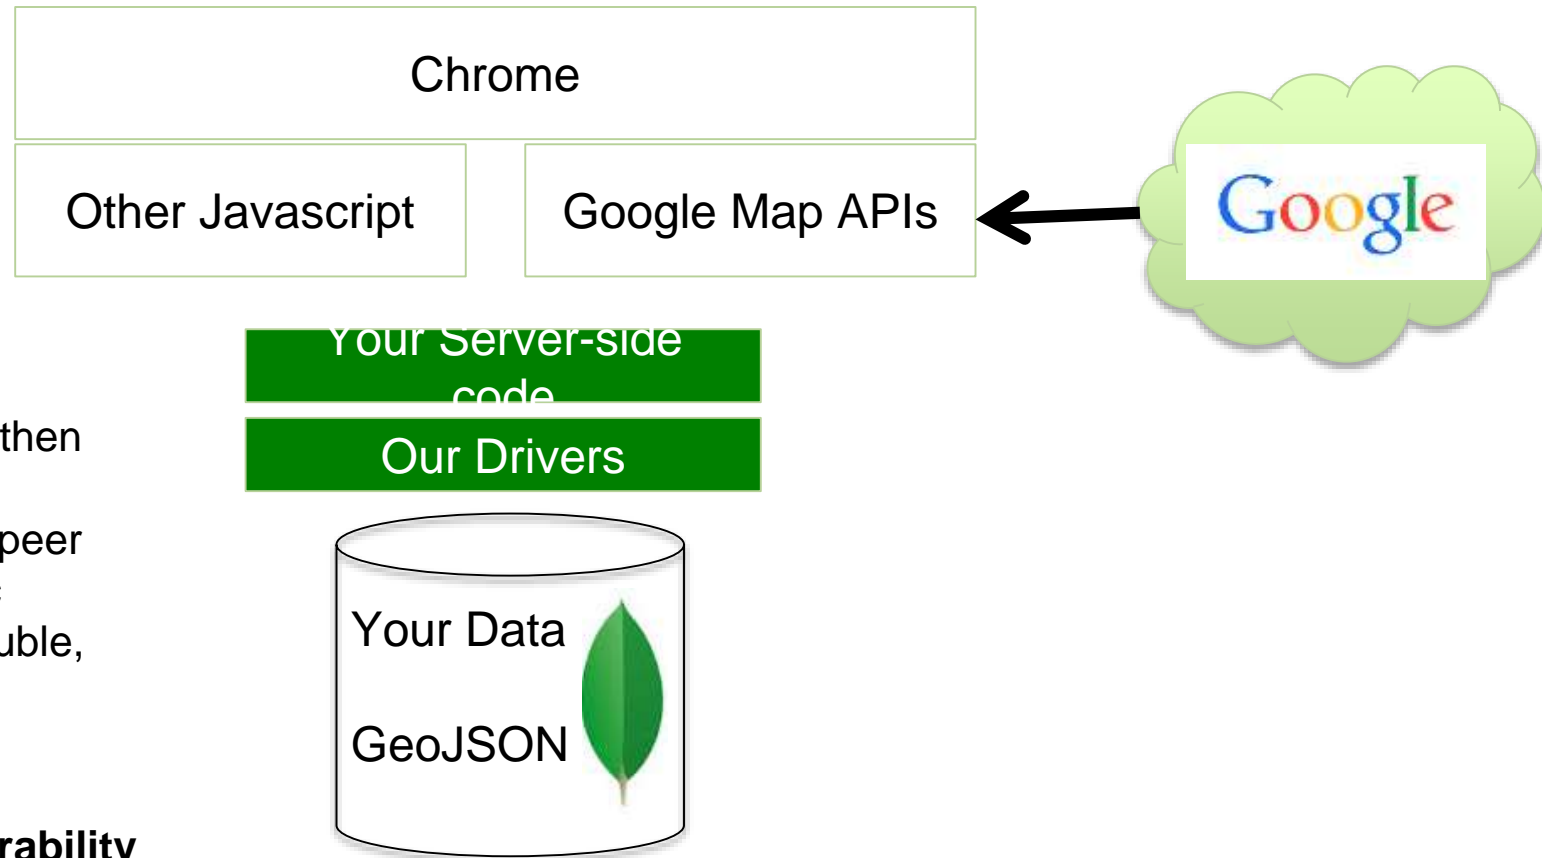
- Geospatial queries do not necessarily require an index
- 2d vs. 2dsphere
- Geo stacks and the Ecosystem
- MongoDB and OpenGIS and OpenGEO
- Google Maps
- MEAN
- A Sampling of Geo Solutions
- Mr. Smiley, etc.
- Integration with esri and shapefiles
- esri shapefile cracking

Clever Hacks

- John Page date thing
- Mr. Smiley
- Wildfire
- Push pts on a LineString and check for intersects
- Perimeter & Area of simple polygon
- makeNgon

MongoDB handles your data + geo

Google handles the maps



- Organization unit is document, then collection
- Geo data can contain arbitrary peer data or higher scope within doc
- Proper handling of int, long, double, and decimal128
- Dates are REAL datetimes
- **Uniform indexability and querability across geo and “regular” data**